

# DOS Taxonomy: Quiz One

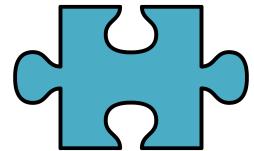
Match the DOS attack classification with its description.

## Attacks:

- 4 Random Scanning
- 2 Permutation Scanning
- 3 Signpost Scanning
- 1 Hitlist Scanning

## Descriptions:

- 1. A portion of a list of targets is supplied to a compromised computer.
- 2. All compromised computers share a common pseudo-random permutation of the IP address space.
- 3. Uses the communication patterns of the compromised computer to find new target.
- 4. Each compromised computer probes random addresses.



## DOS Taxonomy: Quiz Two

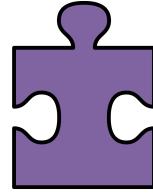
Match the DOS attack classification with its description.

### Attacks:

- 2 Subnet Spoofing
- 1 Random Spoofing
- 3 Fixed Spoofing

### Descriptions:

- 1. Generate 32-bit numbers and stamp packets with them.
- 2. Generate random addresses within a given address space.
- 3. The spoofed address is the address of the target.



# DOS Taxonomy: Quiz Three

Match the DOS attack classification with its description.

Attacks:

- 2 Server Application
- 3 Network Access
- 1 Infrastructure

Descriptions:

- 1. The motivation of this attack is a crucial service of a global internet operation, for example core router
- 2. The attack is targeted to a specific application on a server
- 3. The attack is used to overload or crash the communication mechanism of a network.



## Network DoS:

Goal: take out a large site with little computing

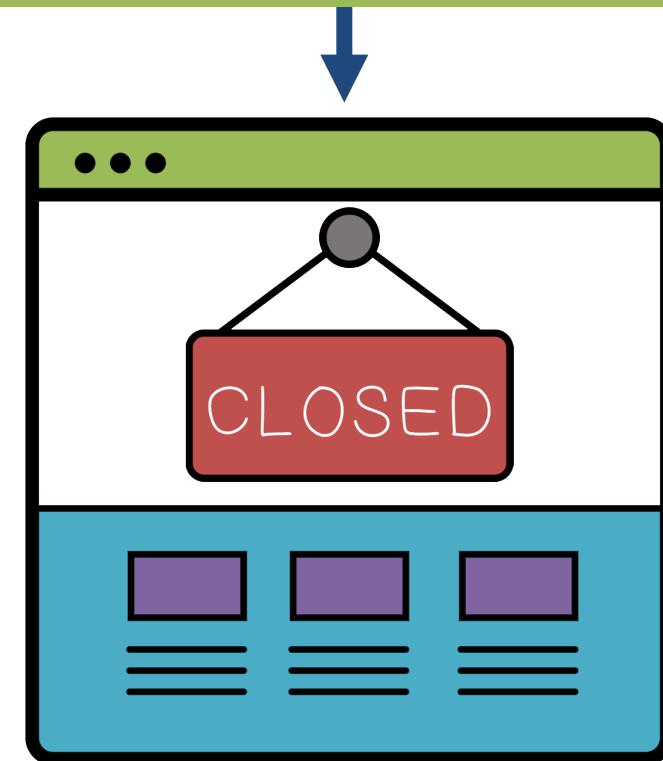


How:

- Amplification
  - Small number of packets



BIG EFFECT





## Network DoS:

Two types of amplification attacks:



DoS bug:

- Design flaw allowing one Machine to disrupt a service



DoS flood:

- Command botnet to Generate flood of requests

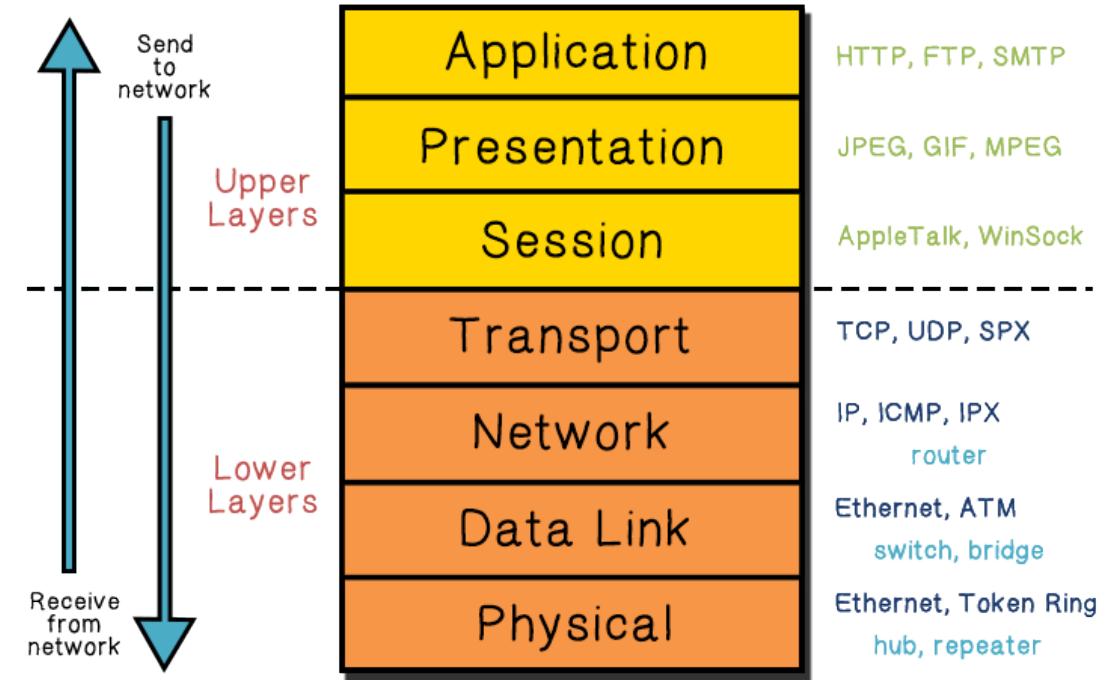


# Network DoS:

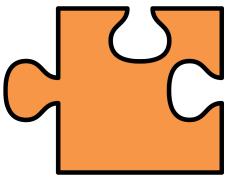
DoS can happen at any layer

Sample DoS at different layers:

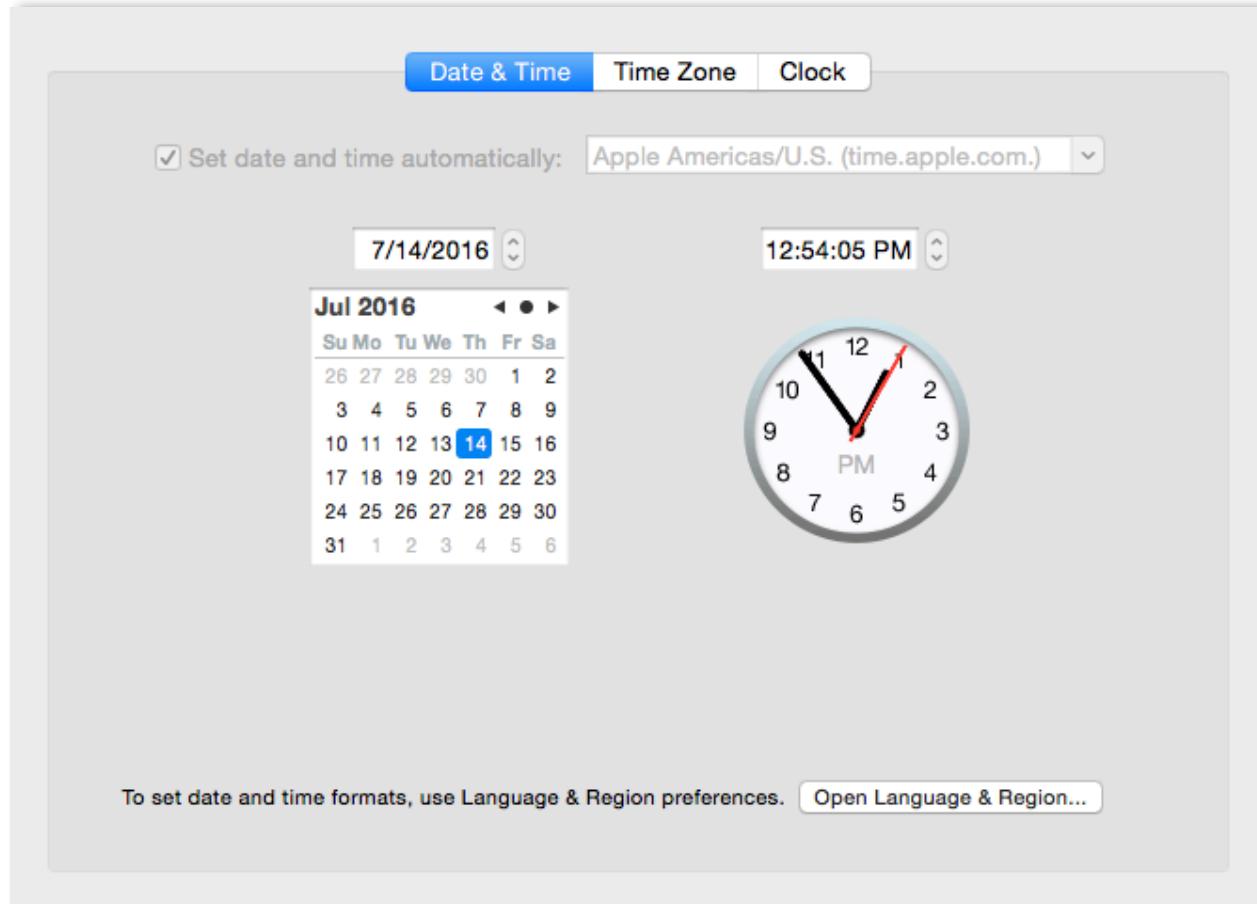
- Link
- TCP/UDP
- Application



Sad truth: Current internet not designed to handle DDoS attacks

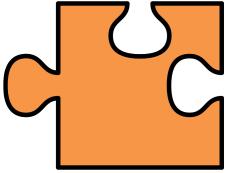


# Amplification Quiz



## NTP – Network Time Protocol

Used to synchronize machines and their clocks.

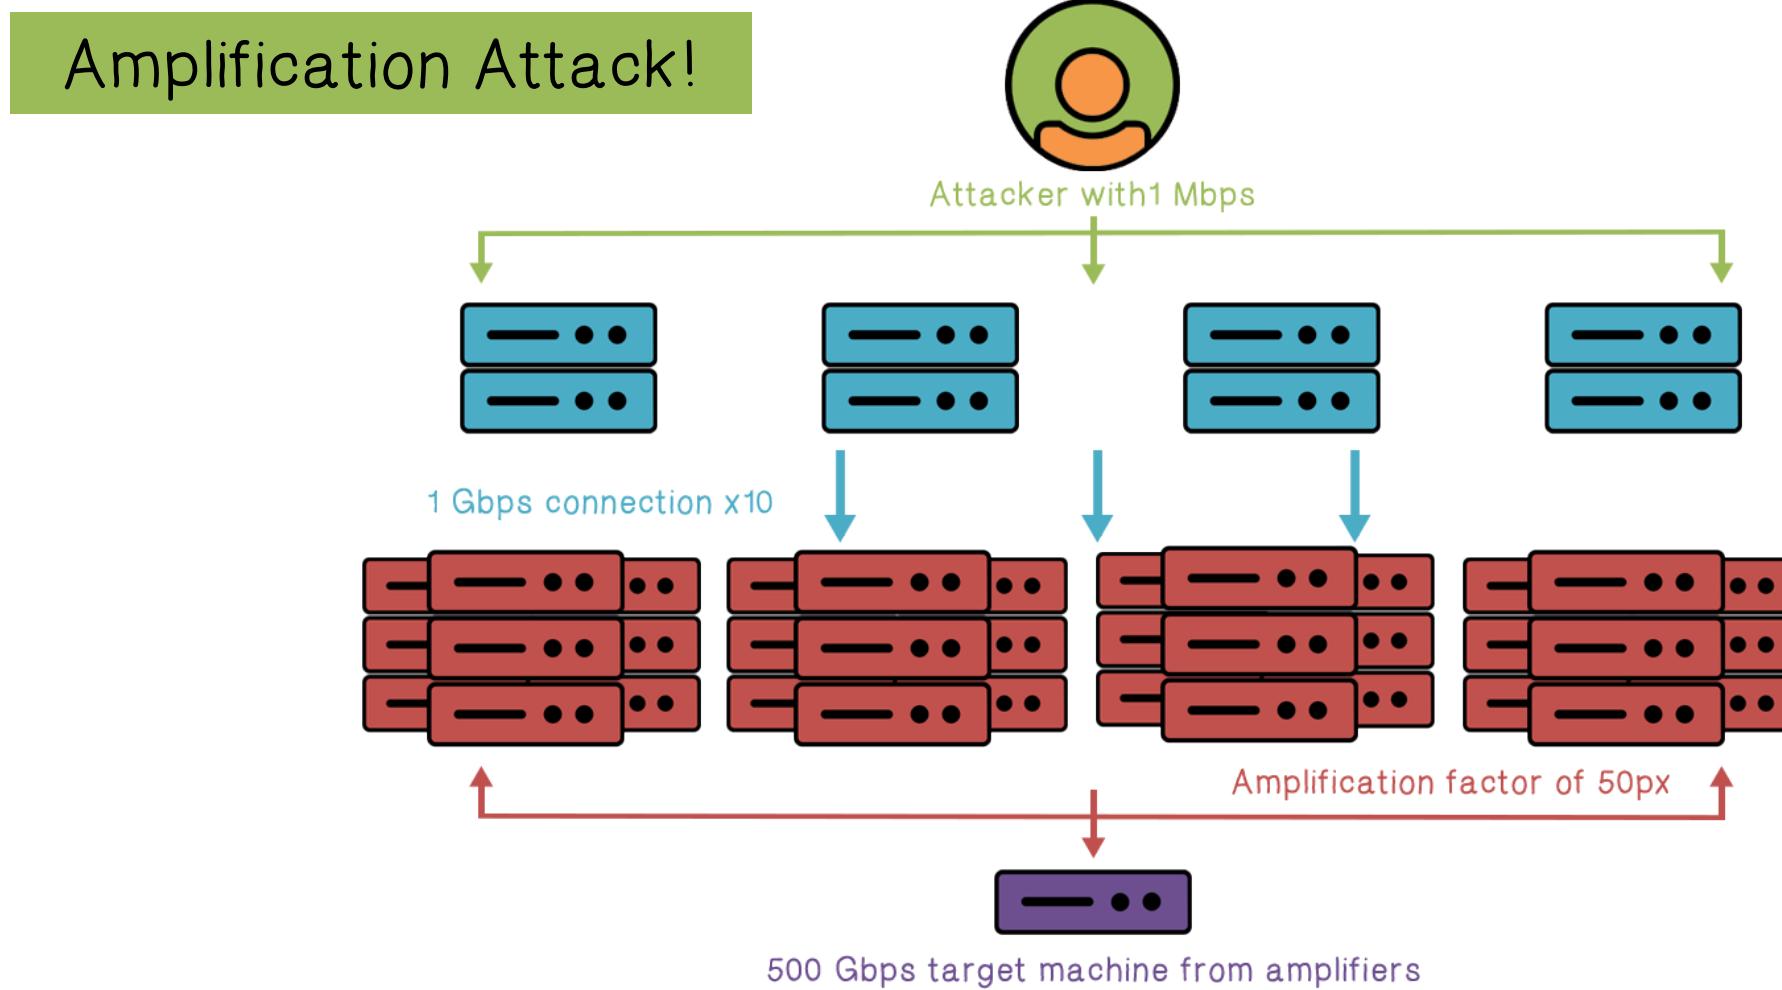


## Amplification Quiz

Which of these are reasons why the UDP-based NTP protocol is particularly vulnerable to amplification attacks?  
Select all that are true.

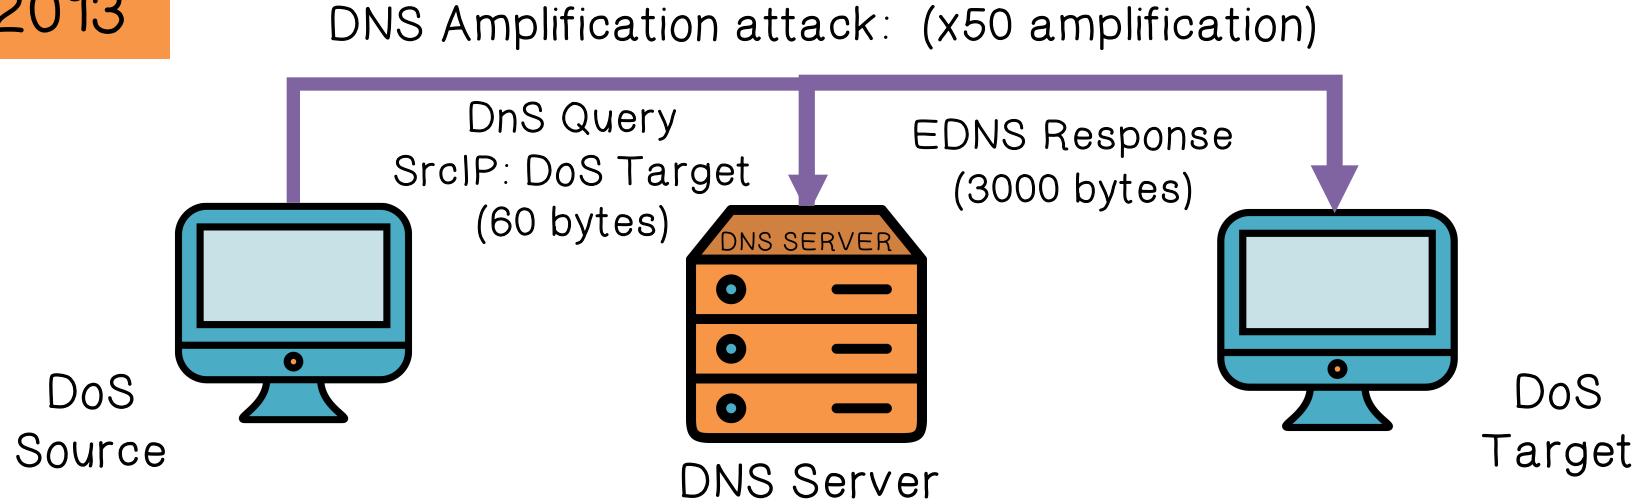
- A small command can generate a large response.
- Vulnerable to source IP spoofing.
- It is difficult to ensure computers communicate only with legitimate NTP servers.

# 🔊 Amplification Example



# 🔊 Amplification Example

March 2013



2006: 0.58M open resolvers on Internet (*Kaminsky-Shiffman*)

2014: 28M open resolvers ([openresolverproject.org](http://openresolverproject.org))

➡ March 2013: DDoS attack generating 309 Gbps for 28 mins

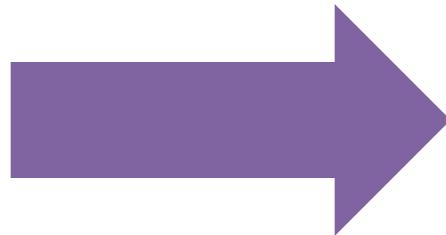


# IP Header Format

Connectionless

Unreliable

Best Effort



Version	Header Length
	Type of Service
	Total Length
	Identification
Flags	Fragment Offset
	Time to Live
	Protocol
	Header Checksum
Source Address of Originating Host	
Destination Address of Target Host	
	Options
	Padding
	IP Data



# TCPHeader Format

# Session Based

# Congestion control

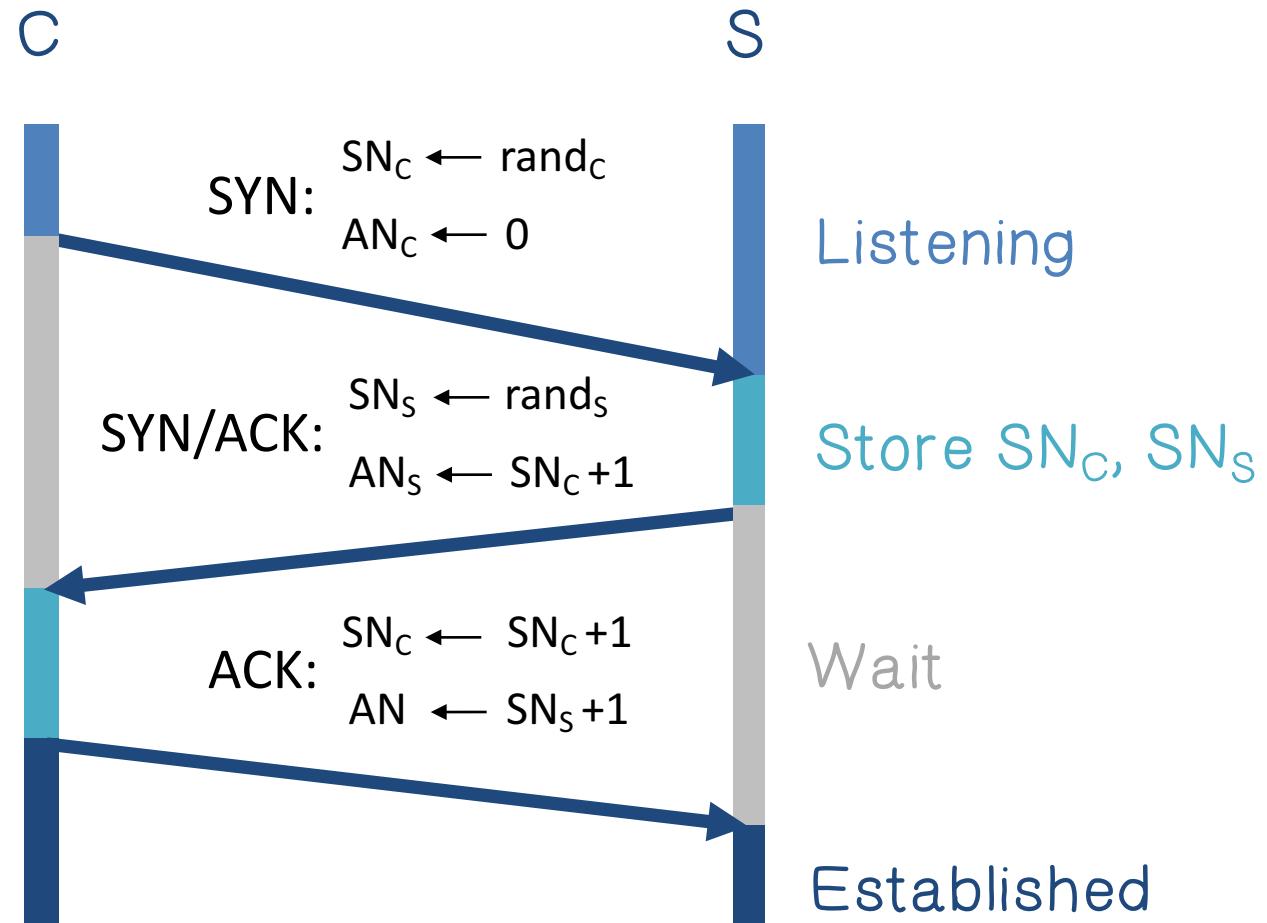
# In order delivery

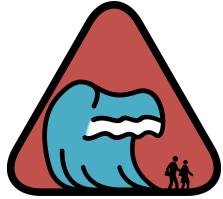


Source Port	Dest Port
SEQ Number	
ACK Number	
G	U
R	A
C	P
K	S
	H
	P
	S
	S
	Y
	N
	F
	I
	N
Other stuff	



# TCP Handshake

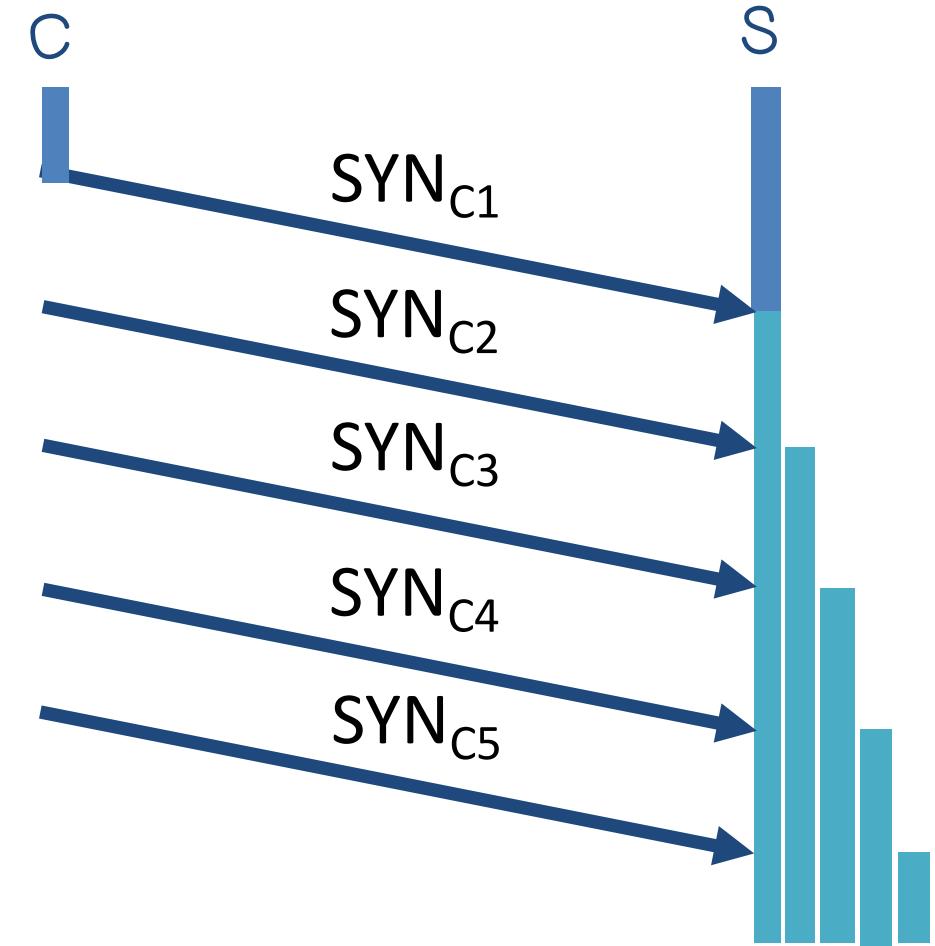


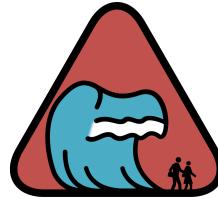


# TCP SYN Flood I: low rate (DoS Bug)

 Single machine:

- SYN Packets with random source IP addresses
- Fills up backlog queue on server
- No further connections possible





# TCP SYN Flood I

A classic SYN flood example



## MS Blaster worm (2003)

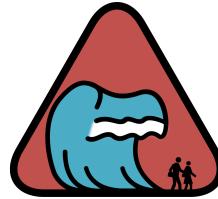
- Infected machines at noon on Aug 16th:
- SYN flood on port 80 to windowsupdate.com
- 50 SYN packets every second
  - each packet is 40 bytes
- Spoofed source IP: a.b.X.Y where X,Y random



## MS Solution



new name:  
windowsupdate.microsoft.com



# TCP SYN Flood I

## Low rate SYN flood defenses



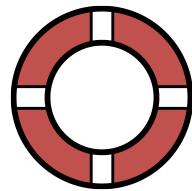
Non-solution:

- Increase backlog queue size or decrease timeout



Correct Solution:

- Syncookies: remove state from server
- Small performance overhead



# SYN COOKIES



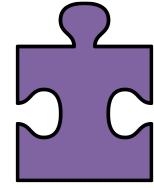
Idea: use secret key and data in packet to generate server SN

Server responds to Client with SYN-ACK cookie:

- T = 5-bit counter incremented every 64 secs.
- L =  $\text{MAC}_{\text{key}}(\text{SAddr}, \text{SPort}, \text{DAddr}, \text{DPort}, \text{SN}_c, \text{T})$  [24 bits]
  - key: picked at random during boot
- $\text{SN}_s = (\text{T} . \text{mss} . \text{L})$  ( $|\text{L}| = 24 \text{ bits}$ )
- Server does not save state

Honest client responds with ACK ( $\text{AN}=\text{SN}_s + 1$ ,  $\text{SN}=\text{SN}_c + 1$ ):

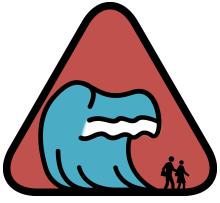
- Server allocates space for socket only if valid  $\text{SN}_s$



## Syn Cookies Quiz

Select all the true statements:

- SYN cookies require modified versions of TCP
- SYN cookies lead to overall slower performance
- The server must reject all TCP options because the server discards the SYN queue entry

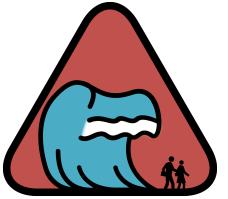


## SYN Floods II: Massive flood

⌚ Command bot army to flood specific target: (DDoS)

- ➡ 20,000 bots can generate 2Gb/sec of SYNs (2003)
- ➡ At web site:
  - Saturates network uplink or network router
  - Random source IP ➡ attack SYNs look the same as real SYNs



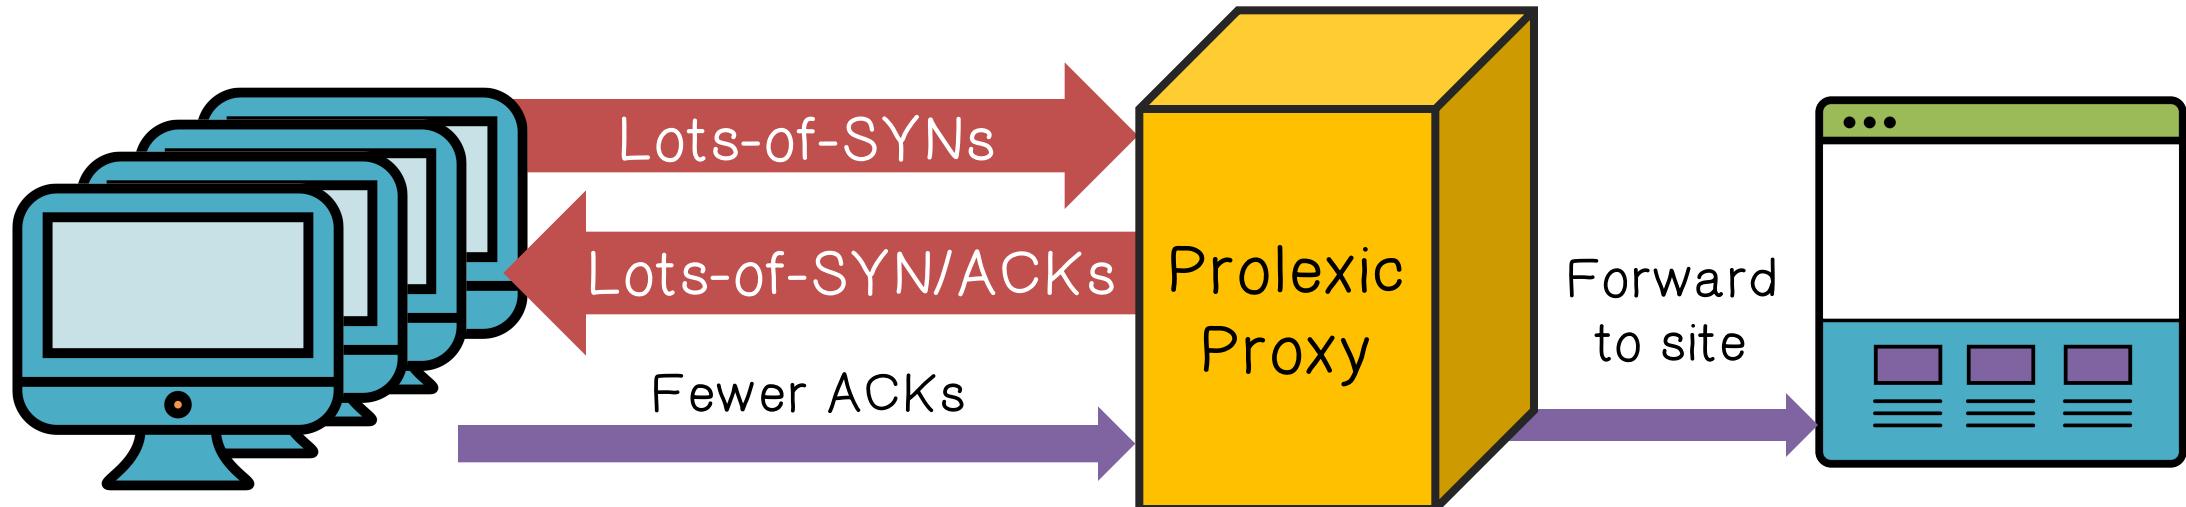


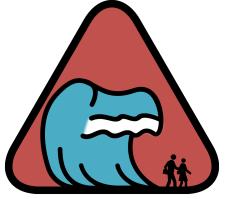
## SYN Floods II: Massive flood

Prolexic / CloudFlare



Idea: only forward established TCP connections to site





# Stronger attacks: TCP connection flood



Command bot army:

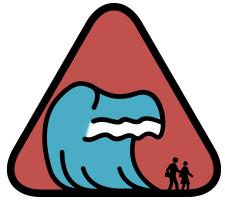
- Complete TCP connection to web site
- Send short HTTP HEAD request
- Repeat

Will bypass SYN flood protection proxy but:

Attacker can no longer use random source IPs

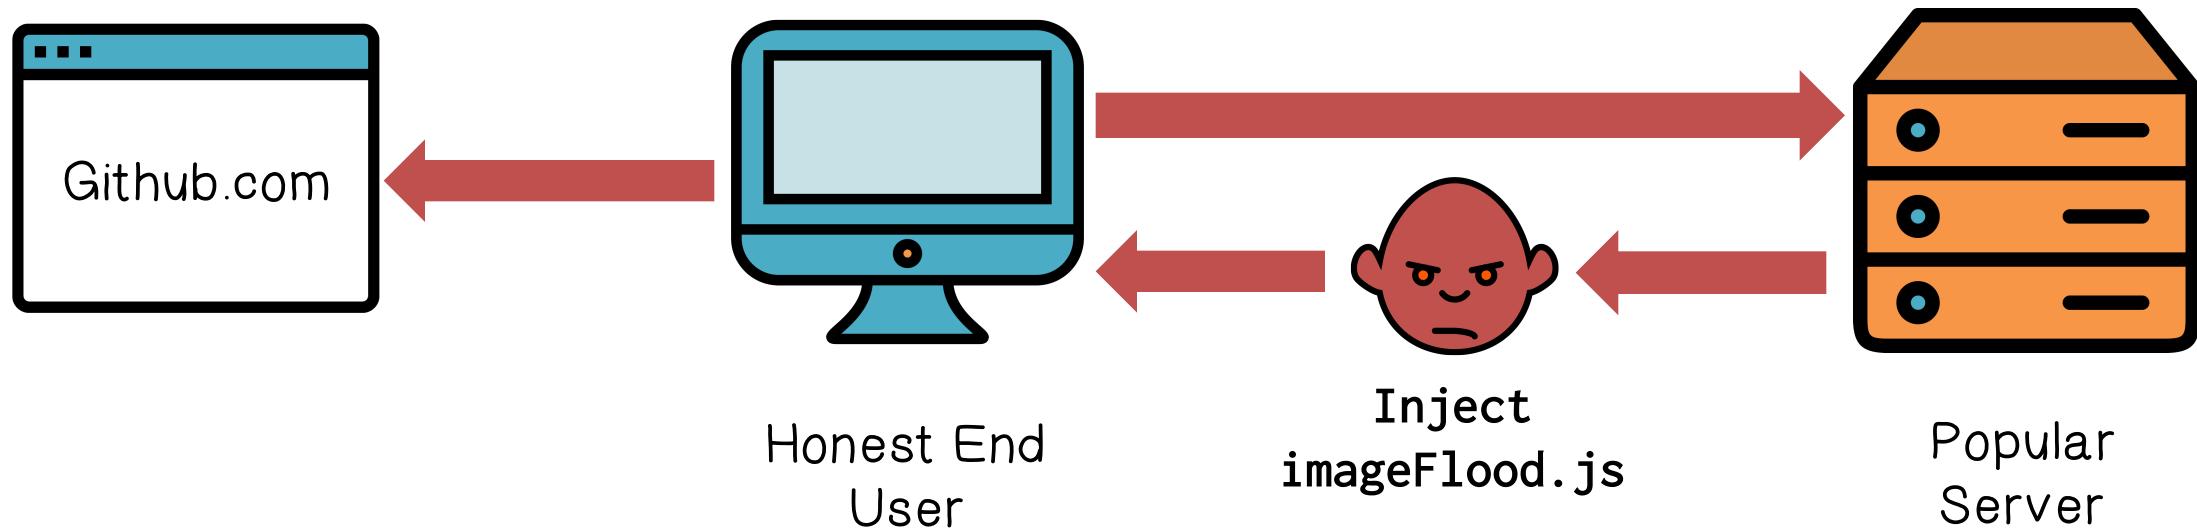
Reveals location of bot zombies

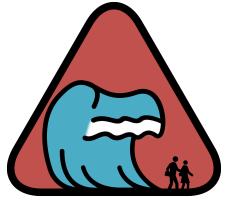
Proxy can now block or rate-limit bots



# A real-world example: GitHub(3/2015)

Javascript-based DDoS:

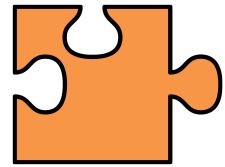




# A real-world example: GitHub(3/2015)

## imageFlood.js

```
Function imgflood() {  
    var TARGET = 'victim-website.com/index.php?'  
    var rand = Math.floor(Math.random() * 1000)  
    var pic = new Image()  
    Pic.src = 'http://'+TARGET+rand+'=val'  
}  
setInterval(imgflood,10)
```

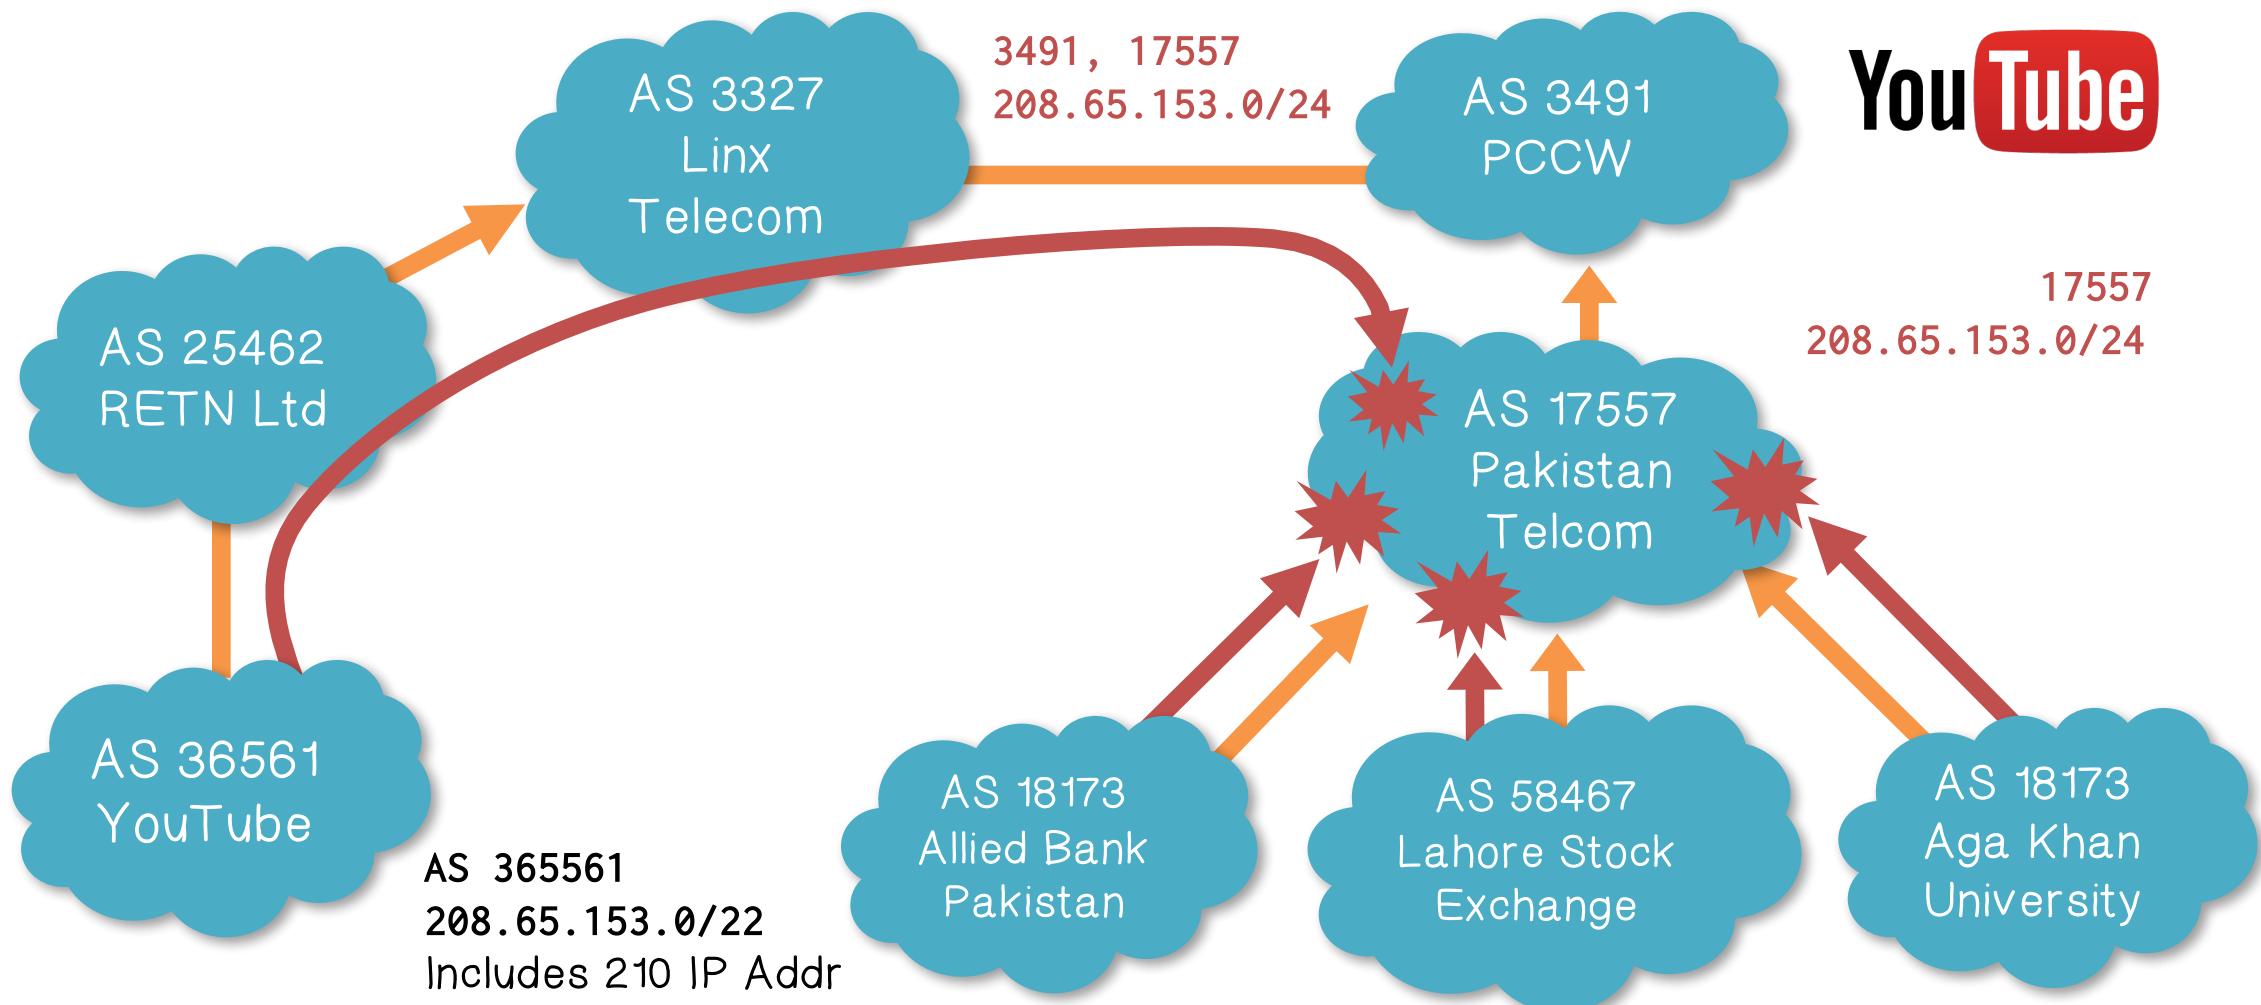


## Flood Attack Quiz

With regards to a UDP flood attack, which of the following statements are true:

- Attackers can spoof the IP address of their UDP packets
- The attack can be mitigated using firewalls
- Firewalls cannot stop a flood because the firewall is susceptible to flooding.

# DoS via route hijacking

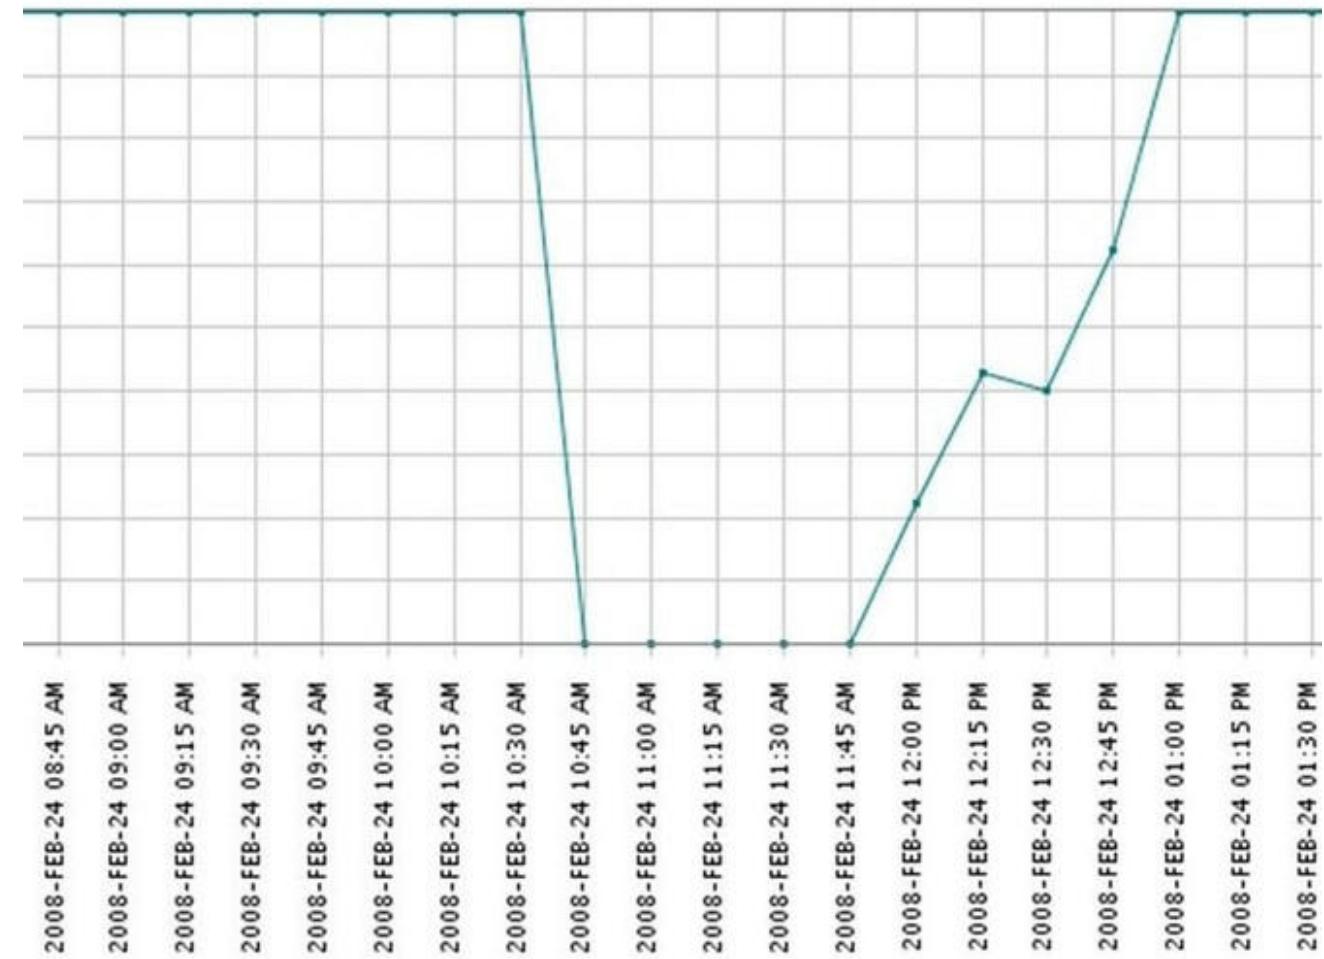


# DoS via route hijacking

DETOUR

You Tube Timeframe:

- 100% at 10:30am
- 0% at 10:45am

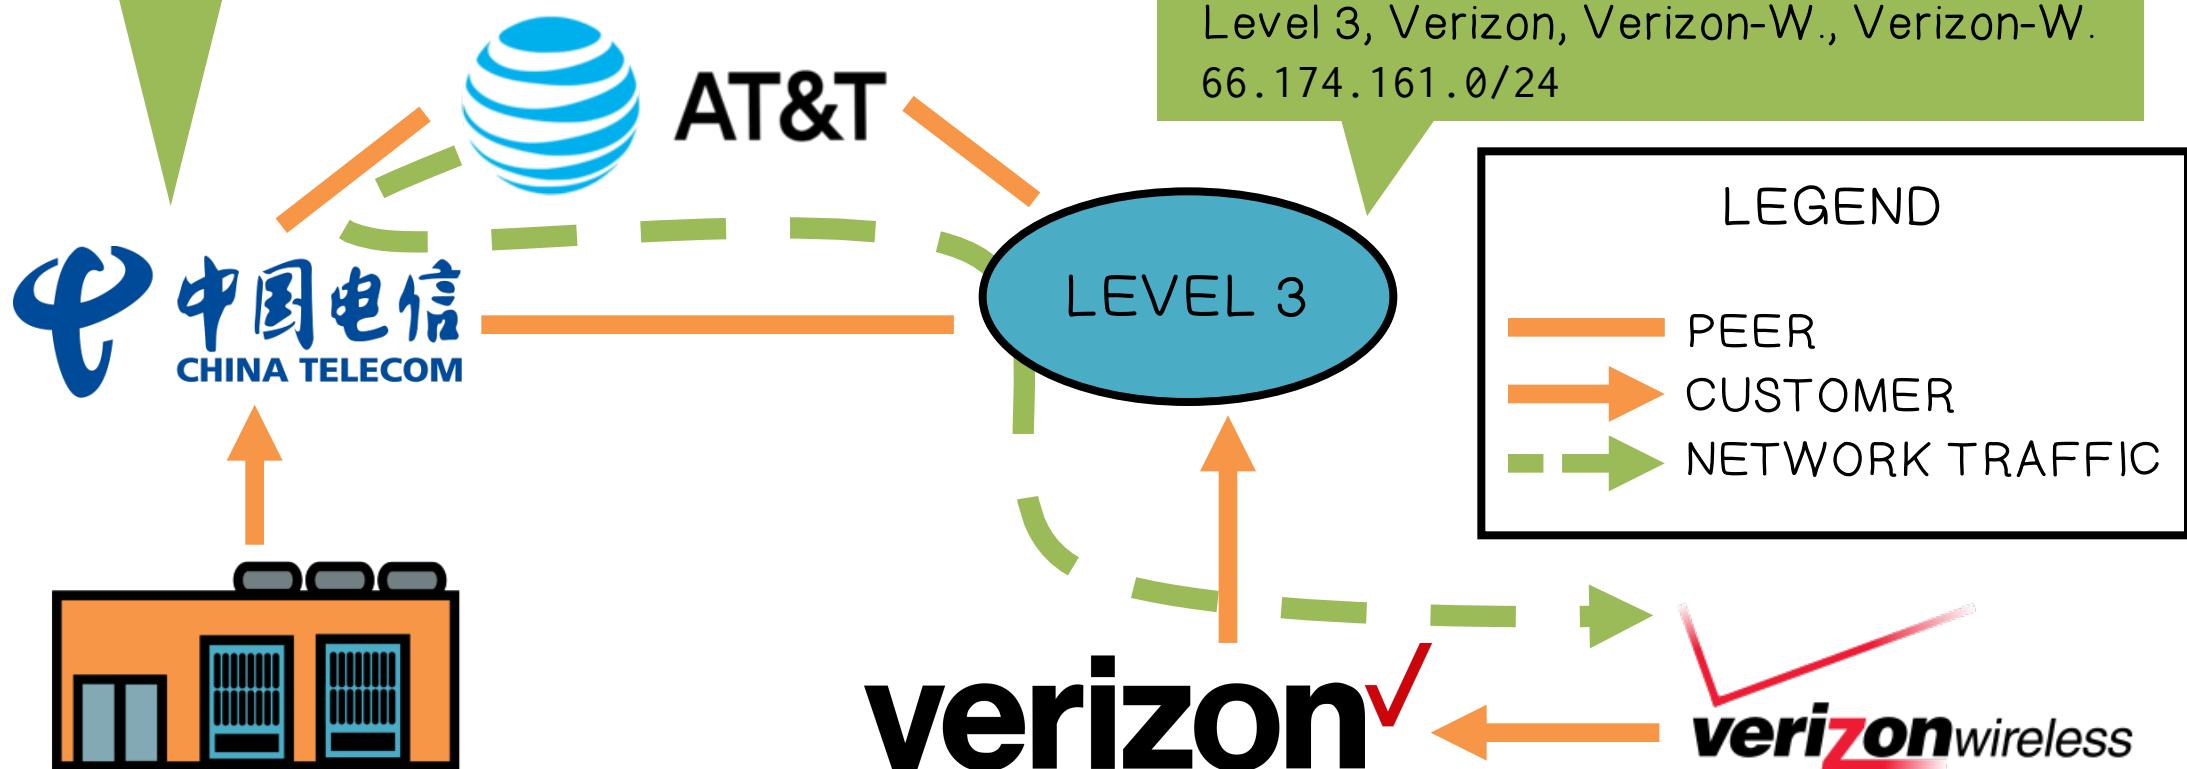


# DoS via route hijacking

DETOUR

China Telecom, China Telecom – DC, China Telecom – DC  
66.174.161.0/24

Level 3, Verizon, Verizon-W., Verizon-W.  
66.174.161.0/24

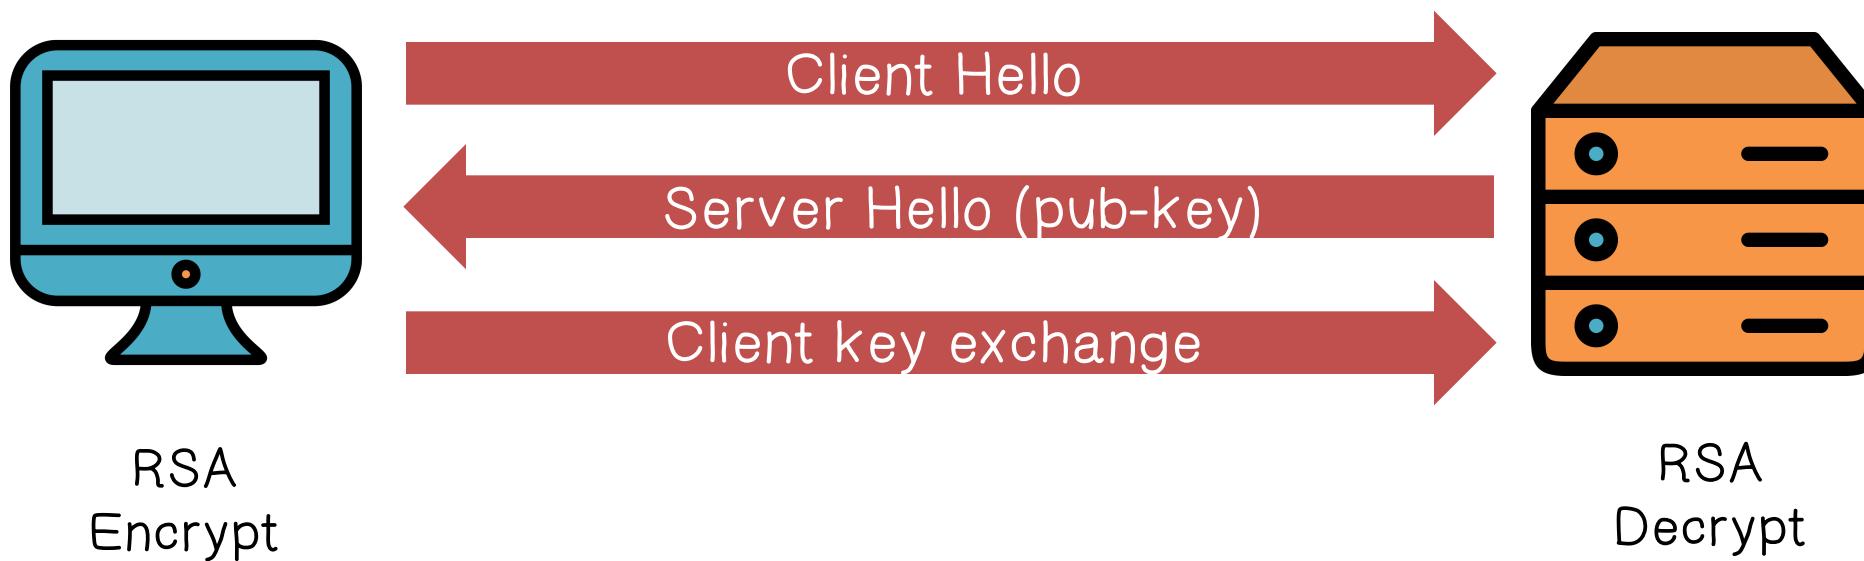




# DoS at Higher Levels



## SSL/TLS handshake [SD'03]





# DoS Mitigation



## Client puzzles



Moderately hard problem:

- Given challenge C find X such that

- $\text{LSB}_n (\text{SHA-1}(C || X)) = 0^n$

- Assumption: takes expected  $2^n$  time to solve

- For  $n=16$  takes about .3sec on 1Ghz machine



Idea: slow down attacker



Main point: checking puzzle solution is easy.



# DoS Mitigation

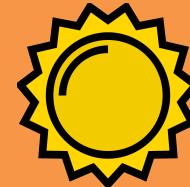


## Client puzzles



During DoS attack:

- Everyone must submit puzzle solution with requests



When no attack:

- Do not require puzzle solution



# DoS Mitigation



## Client puzzles: Examples



### TCP connection floods (RSA '99)

- Example challenge:  $C = \text{TCP server-seq-num}$
- First data packet must contain puzzle solution
  - Otherwise TCP connection is closed





# DoS Mitigation



## Client puzzles: Examples



### SSL handshake DoS: (SD'03)

- Challenge C based on TLS session ID
- Server: check puzzle solution before RSA decrypt



Same for application layer DoS and payment DoS



# DoS Mitigation



## Client puzzles: Benefits and limitations



Hardness of challenge:  $n$

- Decided based on DoS attack volume



Limitations:

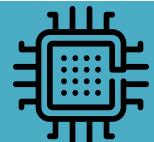
- Requires changes to both clients and servers
- Hurts low power legitimate clients during attack:
  - Clients on cell phones and tablets cannot connect



# DoS Mitigation



Client puzzles: Memory-bound functions



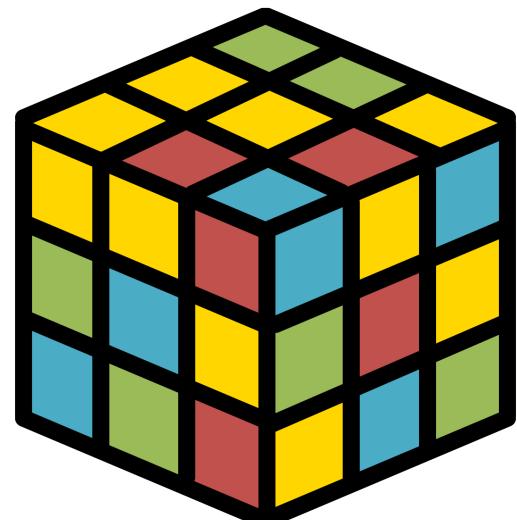
CPU power ratio:

- high end server / low end cell phone = 8000
- Impossible to scale to hard puzzles



Interesting observation:

- Main memory access time ratio:
- high end server / low end cell phone = 2





# DoS Mitigation

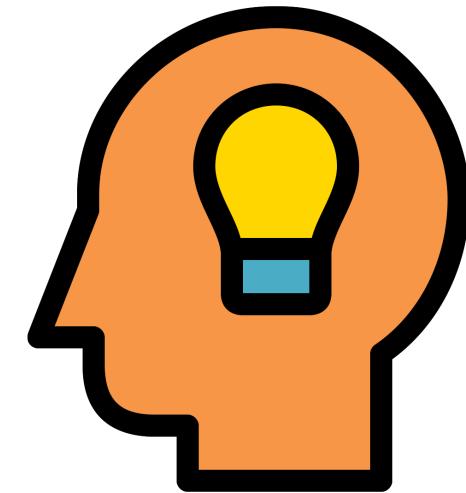


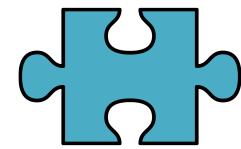
Better puzzles



Solution requires many main memory accesses

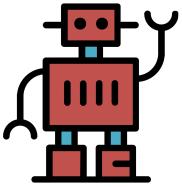
- Dwork-Goldberg-Naor, Crypto '03
- Abadi-Burrows-Manasse-Wobber, ACM ToIT '05



 Puzzle Quiz

Which of the following statements are true?

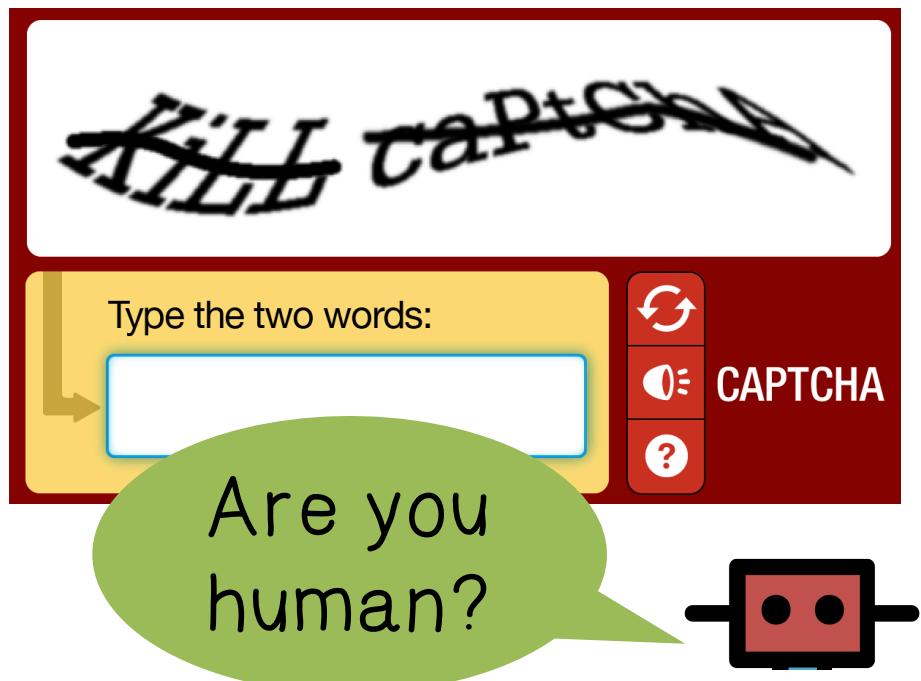
- Client puzzles should be hard to construct. This is an indication of the level of difficulty to solve them.
- Client puzzles should be stateless
- Puzzle complexity should increase as the strength of the attack increases.



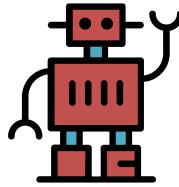
# DoS Mitigation - CAPTCHAs

## CAPTCHA

Completely Automated Public Turing test to tell Computers and Humans Apart



Idea: verify that connection is from a human



# DoS Mitigation - CAPTCHAs

## CAPTCHA

Completely Automated Public Turing test to tell Computers and Humans Apart

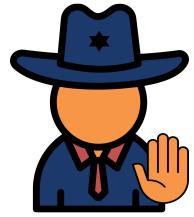


Applies to application layer  
DDoS [Killbots '05]

- During attack: generate CAPTCHAs and process request only if valid solution
- Present one CAPTCHA per source IP address



Idea: verify that connection is from a human

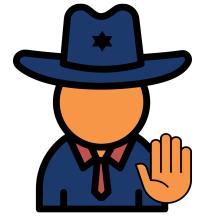


# DoS Mitigation: Source Identification

Goal: identify packet source

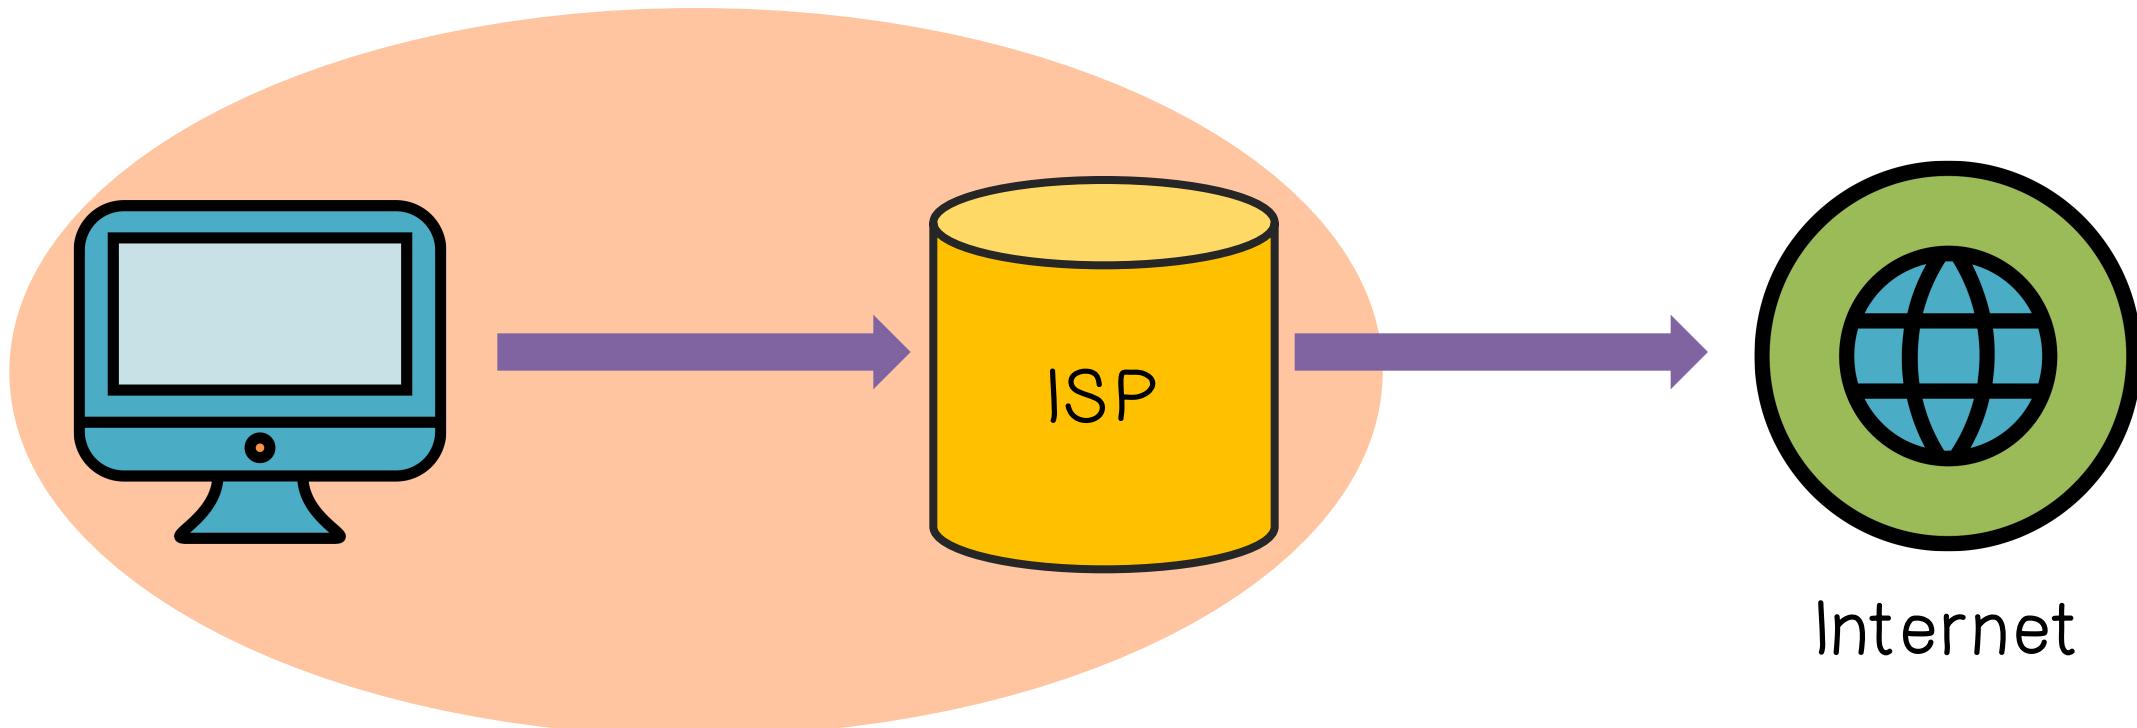


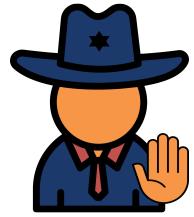
Ultimate goal: block attack at the source



# DoS Mitigation: Source Identification

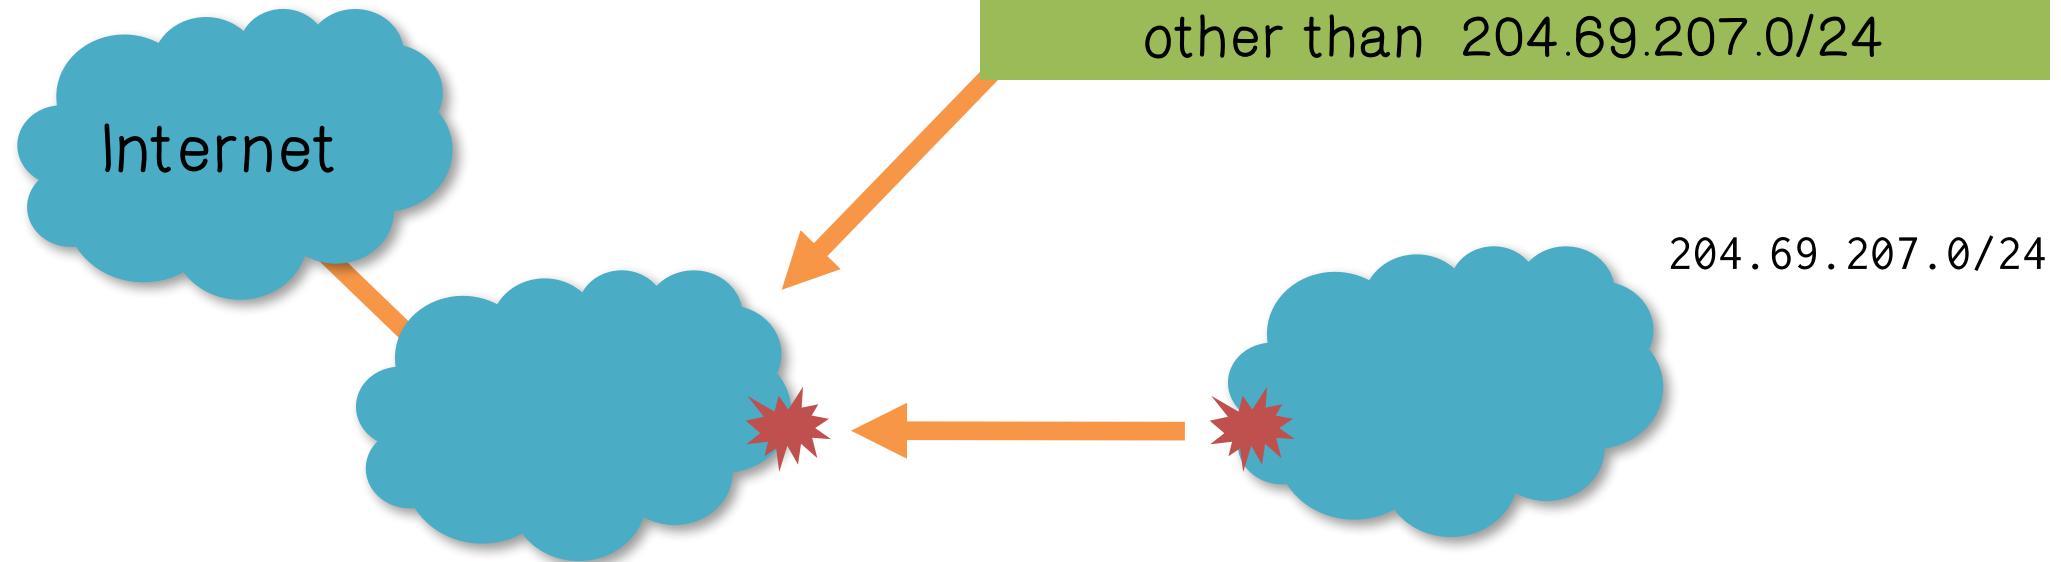
Ingress Filtering



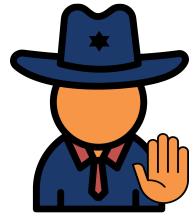


# DoS Mitigation: Source Identification

## Ingress Filtering



Ingress filtering policy: ISP only forwards packets with legitimate source IP



# DoS Mitigation: Source Identification

Ingress Filtering - Implementation problems

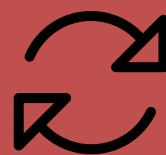


ALL ISPs must do this. Requires global trust.

- If 10% of ISPs do not implement → no defense
- No incentive for deployment

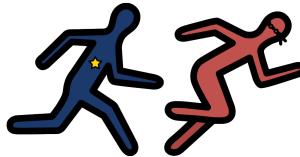


2014:



Recall: 309 Gbps  
attack used only  
3 networks (3/2013)

- 25% of Auto. Systems are fully spoofable ([spoofercmand.org](http://spoofercmand.org))
- 13% of announced IP address space is spoofable



# DoS Mitigation: Traceback

## Traceback [Savage et al. '00]



Goal:

- Given set of attack packets
- Determine path to source

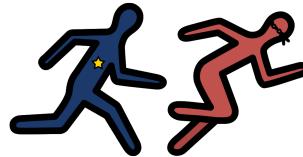


How: change routers to record info in packets



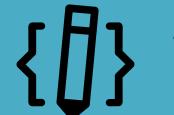
Assumptions:

- Most routers remain uncompromised
- Attacker sends many packets
- Route from attacker to victim remains relatively stable



# DoS Mitigation: Traceback

Simple Method:



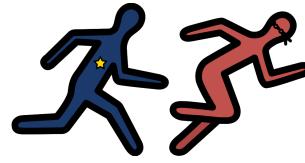
Write path into network packet:

- Each router adds its own IP address to packet
- Victim reads path from packet



Problems:

- Requires space in packet
- Path can be long
- No extra fields in current IP format
  - Changes to packet format too much to expect



# DoS Mitigation: Traceback

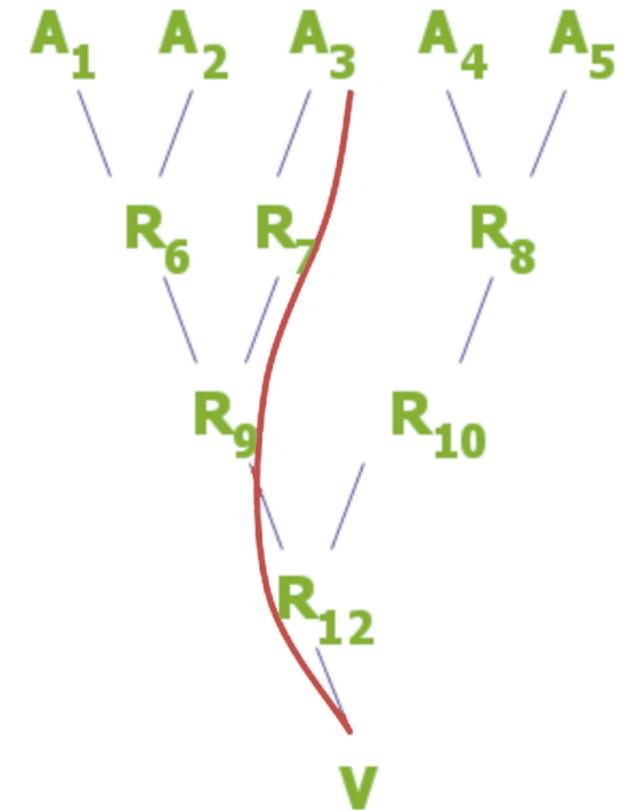


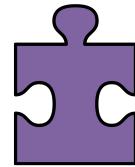
Better Idea

DDoS involves many packets on same path

Store one link in each packet

- Each router probabilistically stores own address
- Fixed space regardless of path length

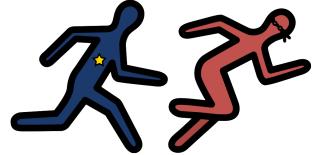




## Traceback Quiz

Which of the following are assumptions that can be made about Traceback?

- Attackers can generate limited types of packets
- Attackers may work alone or in groups
- Attackers are not aware of the tracing mechanism

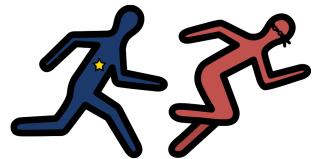


# DoS Mitigation: Edge Sampling



Data fields written to packet:

- Edge: start and end IP addresses
- Distance: number of hops since edge stored



# DoS Mitigation: Edge Sampling

Marking procedure for router R:



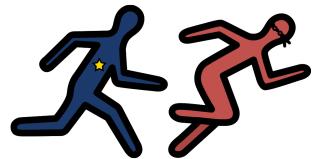
if coin turns up heads (with probability p) then

- write R into start address
- write 0 into distance field



else

- if **distance == 0** write R into end field
- increment distance field

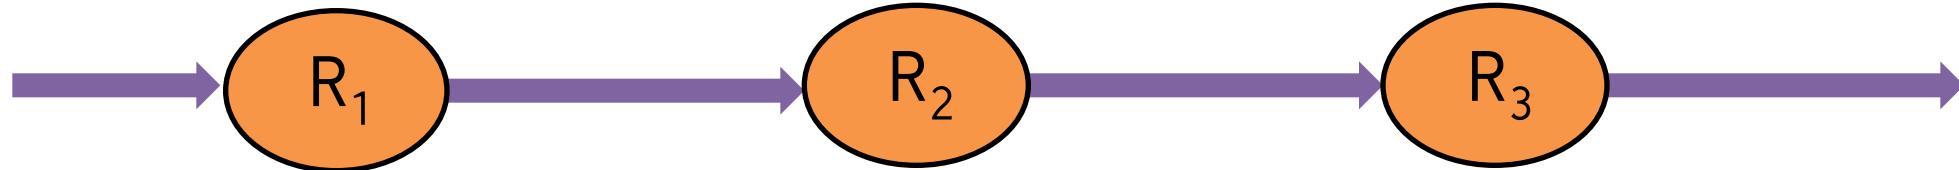
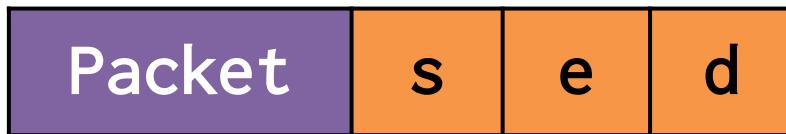


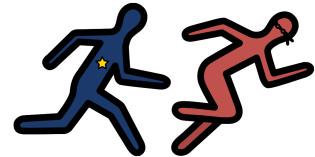
# DoS Mitigation: Edge Sampling



Packet received

- R1 receives packet from source or another router
- Packet contains space for start, end, distance

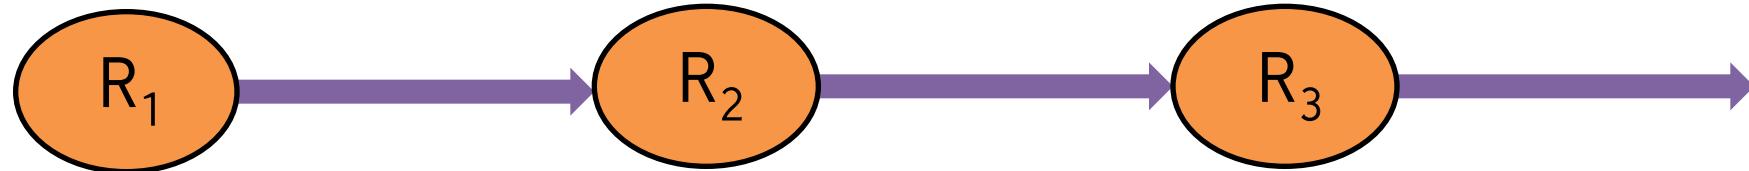
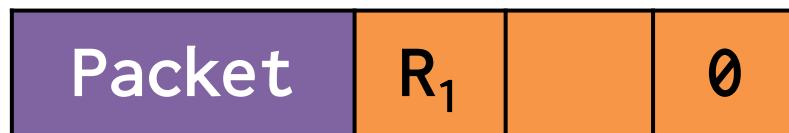


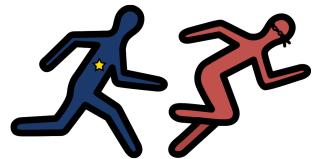


# DoS Mitigation: Edge Sampling

{ } Begin writing edge

- R1 chooses to write start of edge
- Sets distance to 0

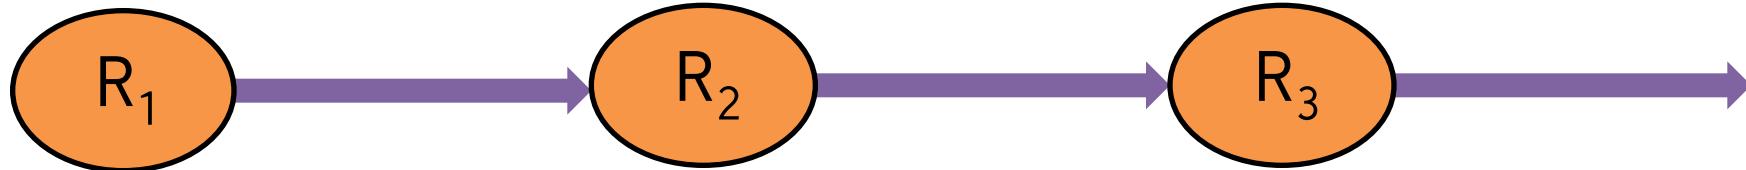
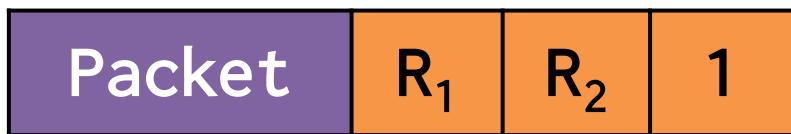


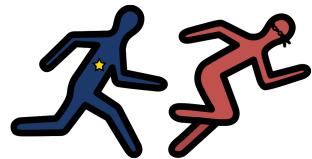


# DoS Mitigation: Edge Sampling

{ } Finish writing edge

- R2 chooses not to overwrite edge
- Distance is 0
  - Write end of edge, increment distance to 1



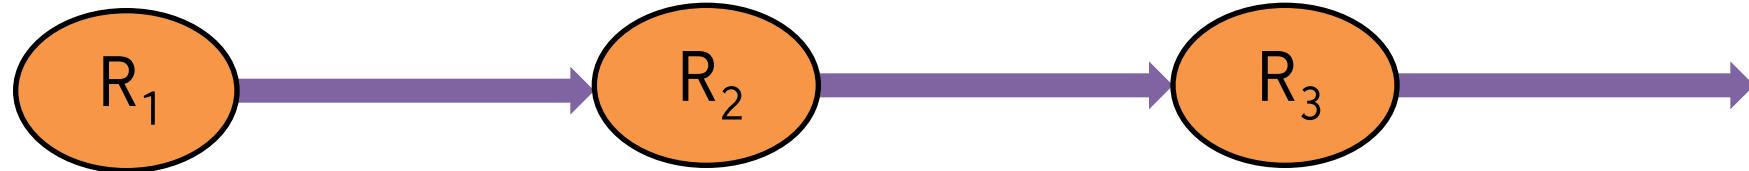
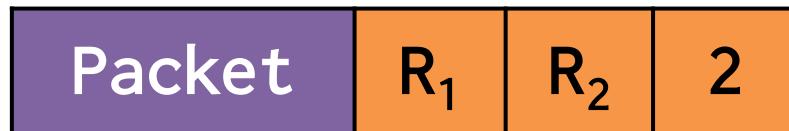


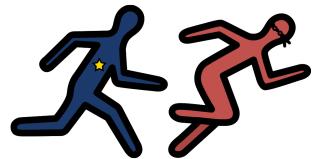
# DoS Mitigation: Edge Sampling



Increment distance

- R<sub>3</sub> chooses not to overwrite edge
- Distance > 0
  - Increment distance to 2





# DoS Mitigation: Edge Sampling

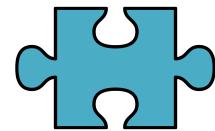


## Path reconstruction

- Extract information from attack packets
- Build graph rooted at victim
  - Each (start,end,distance) tuple provides an edge
- # packets needed to reconstruct path

$$E(X) < \frac{\ln(d)}{p(1-p)^{d-1}}$$

where  $p$  is marking probability,  $d$  is length of path



## Edge Sampling Quiz

Select all the statements that are true for edge sampling:

- Multiple attackers can be identified since edge identifies splits in reverse path
- It is difficult for victims to reconstruct a path to the attacker
- Requires space in the IP packet header



# Reflector Attack [Paxson '01]



3 Victim is flooded by all of the data sent from the DNS Servers

1 Attacker spoofs Victim's IP and sends DNS query to many DNS Servers

2 All DNS servers respond to the DNS query and send data to Victim's IP



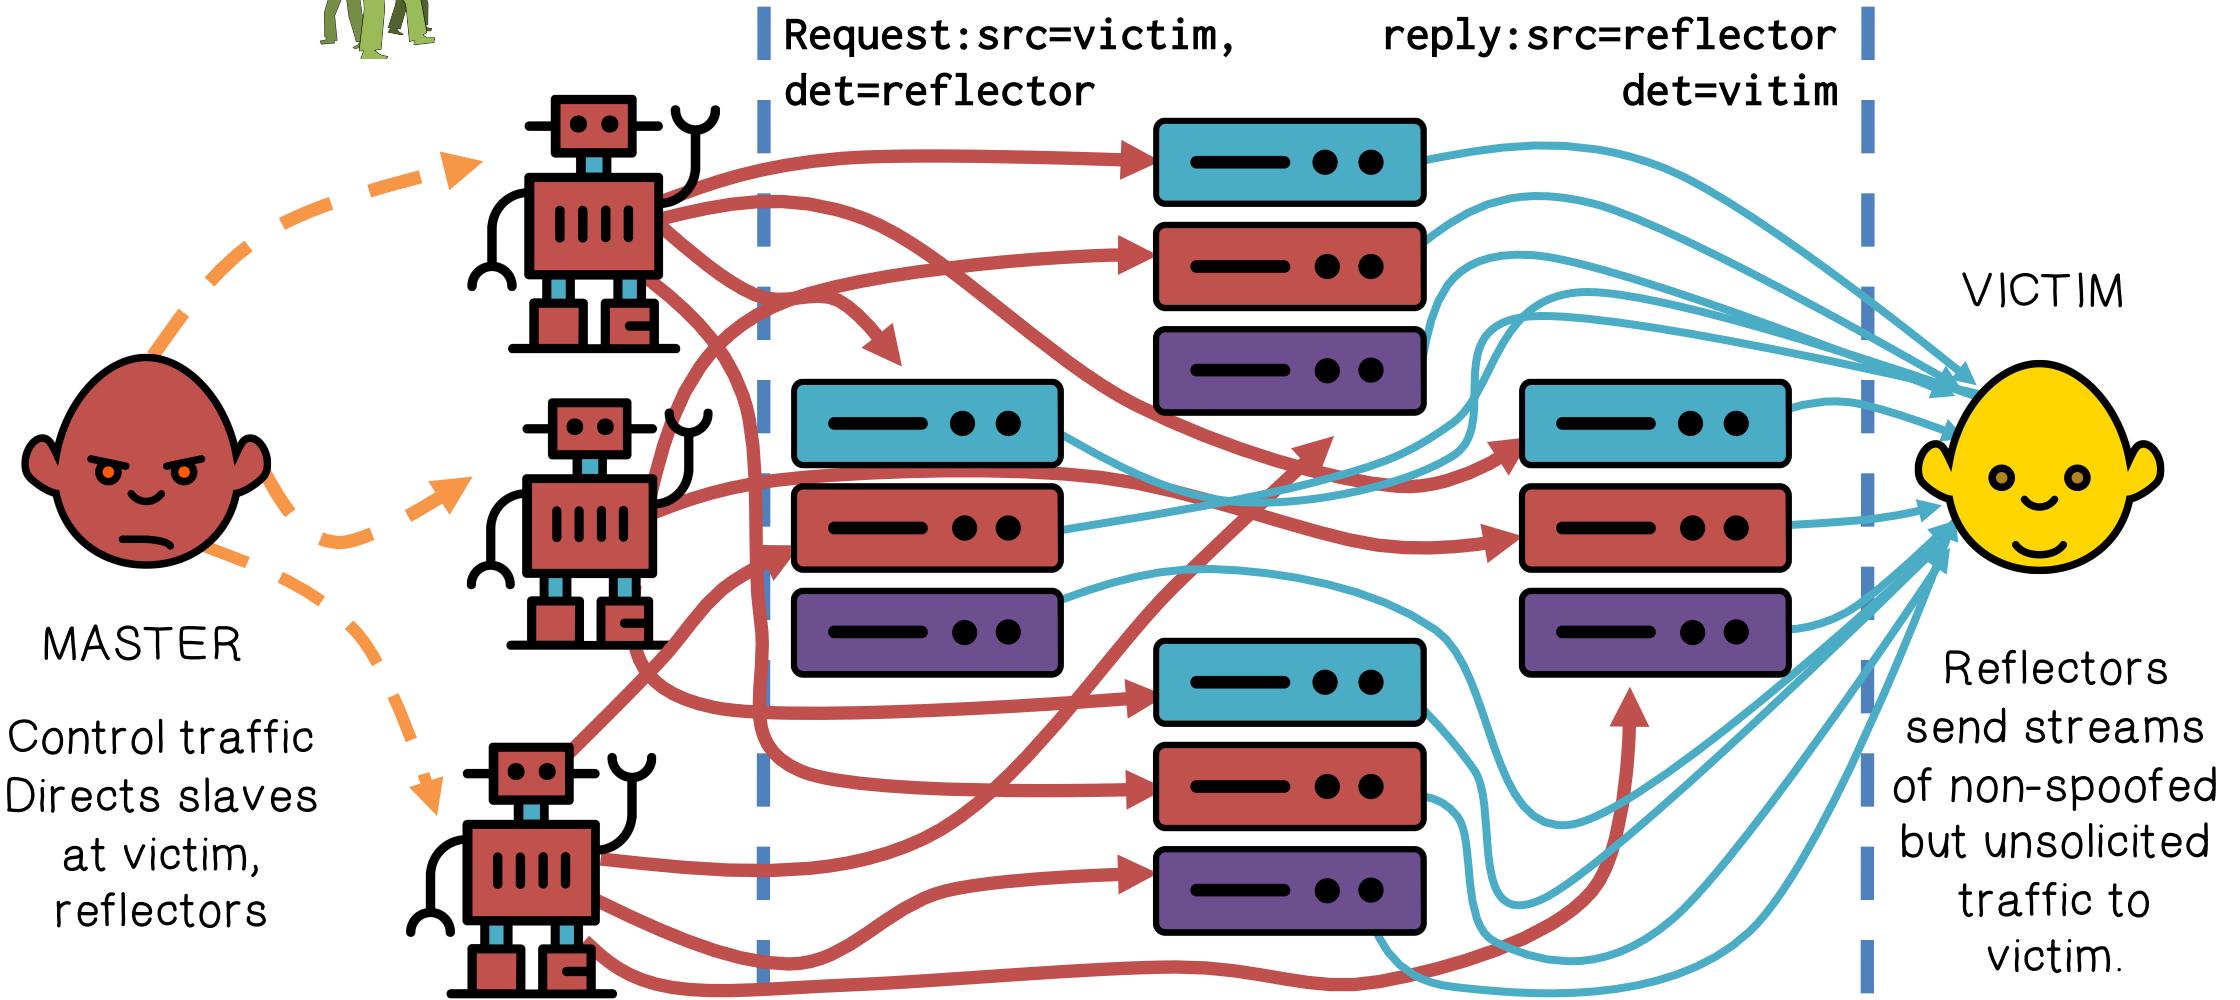
# Reflector Attack [Paxson '01]

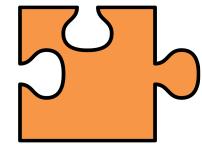
## Examples:

- DNS Resolvers: UDP 53 with victim.com source
  - At victim: DNS response
- Web servers: TCP SYN 80 with victim.com source
  - At victim: TCP SYN ACK packet
- Gnutella servers



# Reflector Attack [Paxson '01]

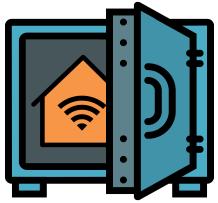




## Reflector Attack Quiz

Self defense against reflector attacks should incorporate:

- Filtering - filter DNS traffic as close to the victim as possible.
- Server redundancy - servers should be located in multiple networks and locations.
- Traffic limiting - traffic from a name server should be limited to reasonable thresholds.



# Capability Based Defense

Anderson, Roscoe, Wetherall

Preventing internet denial-of-service  
with capabilities. SIGCOMM '04.

Yaar, Perrig, and Song

Siff: A stateless internet flow filter to  
mitigate DDoS flooding attacks. IEEE  
S&P '04.

Yang, Wetherall, Anderson

A DoS-limiting network architecture.  
SIGCOMM '05



# Capability Based Defense



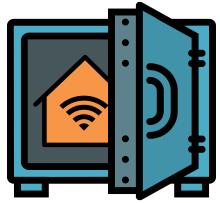
Basic idea:

- ↳ Receivers can specify what packets they want



How:

- Sender requests capability in SYN packet
  - Path identifier used to limit # reqs from one source
- Receiver responds with capability
- Sender includes capability in all future packets



# Capability Based Defense

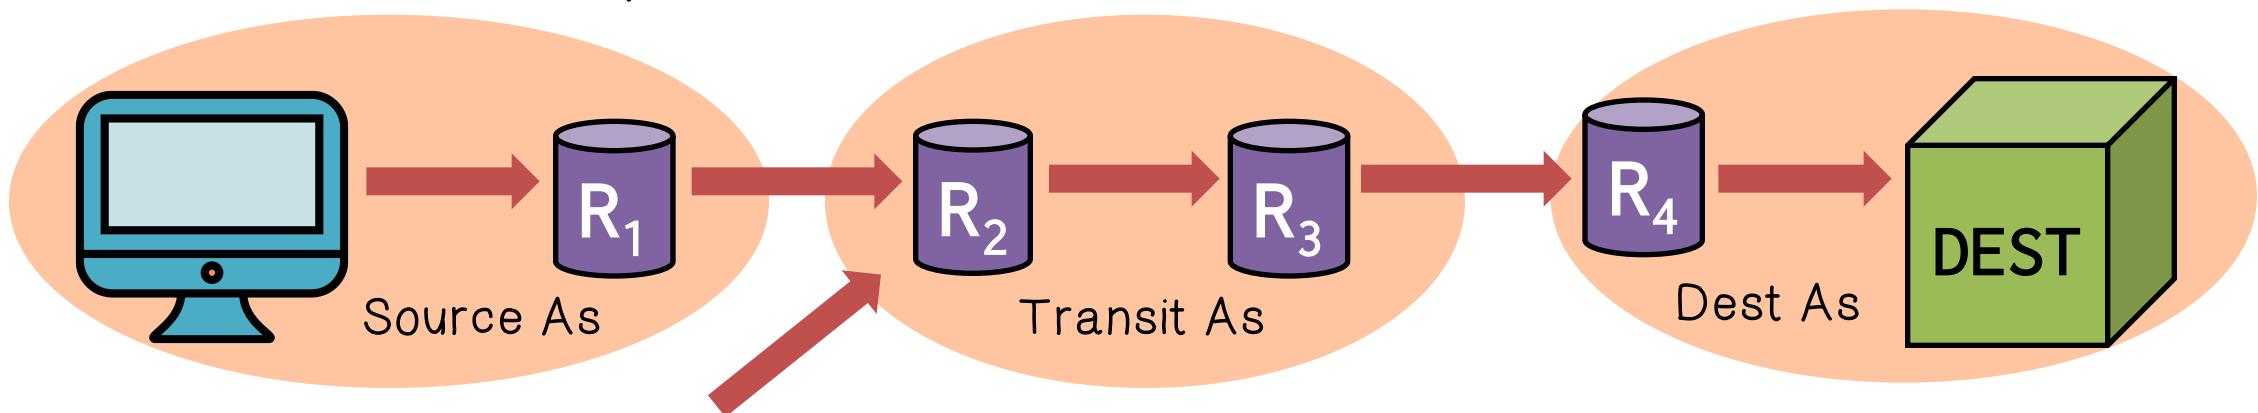


Main point: Routers only forward:

- Request packets, and
- Packets with valid capability

Capabilities can be revoked if source is attacking

- Blocks attack packets close to source





## DoS Summary



Denial of Service attacks are real.



Must be considered at design time.



Sad truth:



Internet is ill-equipped to handle DDoS attacks



Commercial solutions: CloudFlare, Prolexic



Many good proposals for Internet core redesign.